

Invoice Brain

Production Benchmark Report

Real Numbers for Production Readiness Evaluation

Overall Score	68.9/100	NEEDS MAJOR IMPROVEMENTS
Test Date	2026-08-01	
Dataset Size	10,000 invoices	
Chaos Tests	5/5	Passed

1. Executive Summary

This benchmark report presents comprehensive production readiness evaluation results for the Invoice Brain system. The testing methodology encompasses five distinct phases: dataset generation, core metrics benchmarking, load testing under various concurrency scenarios, chaos engineering tests for failure resilience, and cost simulation analysis across multiple volume tiers. The evaluation utilizes a synthetic dataset comprising 100 unique supplier profiles with 100 invoices each, totaling 10,000 invoice documents distributed across 10 distinct format templates including Arabic tax invoices, English commercial invoices, mixed bilingual formats, e-commerce receipts, and specialized industry formats for healthcare, legal services, and restaurant operations.

Key Findings:

- Pattern Hit Rate achieved 92.77% against target of >90%, indicating strong pattern recognition capability
- AI Fallback Rate recorded at 4.56%, well below the 10% threshold, demonstrating effective learning retention
- Collision Rate measured at 0.000%, exceeding expectations with near-zero fingerprint conflicts
- False Match Rate registered at 1.900%, exceeding the 0.5% target and requiring attention
- Processing throughput reached 86537.25 invoices/second under sustained load conditions

2. Core Metrics Benchmark Results

The core metrics phase evaluates fundamental system performance indicators that directly impact production viability. Each metric is measured against predefined thresholds established based on enterprise deployment requirements and industry standards for intelligent document processing systems. The benchmark processes all 10,000 invoices through the complete extraction pipeline including layout-based fingerprinting, pattern matching, confidence scoring, and result verification stages.

Primary Performance Indicators:

Metric	Target	Actual Result	Status
--------	--------	---------------	--------

Pattern Hit Rate	>90%	92.77%	PASS
AI Fallback Rate	<10%	4.56%	PASS
Collision Rate	<0.1%	0.000%	PASS
False Match Rate	<0.5%	1.900%	FAIL
Avg Processing Time	<100ms	0.010ms	PASS

Supporting Performance Metrics:

Metric	Value
Throughput	86537.25 invoices/sec
P50 Latency	0.01 ms
P95 Latency	0.01 ms
P99 Latency	0.03 ms
Unique Fingerprints	179
Total Processed	10,000
Cache Hit Rate	0.00%

3. Load Testing Results

Load testing evaluates system behavior under concurrent user scenarios ranging from small business workloads to peak enterprise-scale operations. Each scenario simulates multiple users simultaneously submitting batch invoice processing requests, measuring throughput capacity, latency characteristics, and resource utilization patterns. These tests are critical for validating that the architecture can handle real-world production traffic without degradation in service quality or unacceptable response times.

Scenario: Small Business

Parameter	Value
Concurrent Users	100
Invoices Per User	100
Total Invoices Processed	10,000
Total Processing Time	0.07s
Throughput	146229.01 inv/s
Average Latency	0.01ms
P95 Latency	0.01ms
P99 Latency	0.02ms
Peak Memory Usage	47.0 MB

Scenario: Medium Enterprise

Parameter	Value
Concurrent Users	500
Invoices Per User	100
Total Invoices Processed	50,000
Total Processing Time	0.20s

Throughput	248533.02 inv/s
Average Latency	0.00ms
P95 Latency	0.01ms
P99 Latency	0.01ms
Peak Memory Usage	41.4 MB

Scenario: Large Enterprise

Parameter	Value
Concurrent Users	1000
Invoices Per User	100
Total Invoices Processed	100,000
Total Processing Time	0.33s
Throughput	303227.39 inv/s
Average Latency	0.00ms
P95 Latency	0.01ms
P99 Latency	0.01ms
Peak Memory Usage	68.7 MB

Scenario: Peak Load

Parameter	Value
Concurrent Users	5000
Invoices Per User	20
Total Invoices Processed	100,000
Total Processing Time	0.38s

Throughput	264318.90 inv/s
Average Latency	0.00ms
P95 Latency	0.01ms
P99 Latency	0.01ms
Peak Memory Usage	61.2 MB

4. Chaos Engineering Test Results

Chaos testing validates system resilience under failure conditions by intentionally disrupting critical dependencies and infrastructure components. This methodology ensures graceful degradation rather than catastrophic failure when inevitable production incidents occur. The test suite covers cache layer failures, AI service unavailability, database performance degradation, malformed input handling, and memory pressure scenarios. A robust system should maintain acceptable service levels even when individual components fail or operate outside normal parameters.

Test Scenario	Result	Summary
Redis Down (Cache Miss Storm)	PASSED	Processed 1000 invoices without cache in 0.01
AI Provider Failure	PASSED	96.0% of requests survived via pattern fallba
PostgreSQL Slow	PASSED	Slow DB added 25000ms overhead, system remain
Empty/Malformed Input	PASSED	6/6 malformed inputs handled gracefully
Memory Pressure	PASSED	Processed 10K large invoices, memory increase

Chaos Test Summary: 5/5 tests passed (100% success rate)

5. Cost Simulation Analysis

Cost simulation projects operational expenses across three volume tiers representing different business scales: mid-volume operations at 100,000 daily invoices, high-volume platforms at one million daily invoices, and enterprise-scale processing at ten million monthly invoices. The analysis accounts for the distribution between pattern-based extractions (minimal cost), AI-assisted processing (API costs), and learning-mode operations (initial training overhead). Savings calculations compare hybrid approach costs against pure-AI baseline to quantify the economic value of the learned pattern system.

Volume Tier: 100K/day

Cost Component	Value
Daily Volume	100,000 invoices
Pattern-Based Extractions	92,770 (92.8%)
AI-Processed Invoices	4,560 (4.6%)
Learning Mode Invoices	2,670
AI Cost (Daily)	\$9.12
AI Cost (Monthly)	\$273.60
Savings vs Pure AI	\$190.88/day (95.4%)
Cost Per Invoice	\$0.000091
Total Monthly Estimate	\$773.60

Volume Tier: 1M/day

Cost Component	Value
Daily Volume	1,000,000 invoices
Pattern-Based Extractions	927,700 (92.8%)

AI-Processed Invoices	45,600 (4.6%)
Learning Mode Invoices	26,700
AI Cost (Daily)	\$91.20
AI Cost (Monthly)	\$2736.00
Savings vs Pure AI	\$1908.80/day (95.4%)
Cost Per Invoice	\$0.000091
Total Monthly Estimate	\$3236.00

Volume Tier: 10M/month

Cost Component	Value
Daily Volume	333,333.3333333333 invoices
Pattern-Based Extractions	309,233 (92.8%)
AI-Processed Invoices	15,200 (4.6%)
Learning Mode Invoices	8,900
AI Cost (Daily)	\$30.40
AI Cost (Monthly)	\$912.00
Savings vs Pure AI	\$636.27/day (95.4%)
Cost Per Invoice	\$0.000091
Total Monthly Estimate	\$1412.00

6. Final Assessment & Recommendations

Overall Production Readiness Score: 68.9/100

Assessment Level: NEEDS MAJOR IMPROVEMENTS

Criteria Breakdown:

Criterion	Weight	Target	Actual	Score
Pattern Hit Rate	25%	>90.0%	92.77%	100%
AI Fallback Rate	20%	<10.0%	4.56%	100%
Collision Rate	20%	<0.1%	0.00%	100%
False Match Rate	15%	<0.5%	1.90%	26%
Avg Processing Time	20%	>100.0%	0.01%	0%

Recommendations:

1. Prioritize resolution of False Match Rate exceeding 0.5% threshold significantly.
2. Review and enhance pattern matching algorithm for edge case handling.
3. Implement stronger validation layers before marking extractions as successful.
4. Conduct additional testing with real-world invoice samples before production consideration.
5. Establish clear criteria and timeline for addressing each failed metric.

Identified Risks & Mitigation:

Risk Factor	Severity	Mitigation Strategy
False Match Risk	High	Add field-level validation and cross-reference checks
Pattern Drift	Medium	Implement automated drift detection and versioning
OCR Quality Dependency	Medium	Enhance preprocessing for noisy inputs

Cache Stampede	Low	Implement request coalescing and stale-while-revalidate
----------------	-----	---